

# 03 Dynamic Representation

*Advanced Machine Learning*


Malte Schilling, Neuroinformatics Group, Bielefeld University

# Goals for Today



# Support Vector Machines

# Support Vector Machine



*from Wikipedia:Support-vector machine*

Distance (geometric margin) of data points to border:

$$y^{(i)} \left( \left( \frac{\mathbf{w}}{\|\mathbf{w}\|} \right)^T x^{(i)} + \frac{b}{\|\mathbf{w}\|} \right) = \gamma^{(i)}$$

For details and derivation see (Ng 2018)

# Finding the Largest Margin

$$\gamma = \min_{i=1,\dots,m} \gamma^{(i)} = \min_{i=1,\dots,m} y^{(i)} \left( \left( \frac{\mathbf{w}}{\|\mathbf{w}\|} \right)^T x^{(i)} + \frac{b}{\|\mathbf{w}\|} \right)$$
$$\Rightarrow \max_{\gamma, \mathbf{w}, b} y^{(i)} \left( \mathbf{w}^T x^{(i)} + b \right) \geq \gamma, i = 1, \dots, m, \text{ with } \|\mathbf{w}\| = 1$$

Unfortunately, the constraint on  $\mathbf{w}$  makes this none convex and not directly solvable.

$$\Rightarrow \max_{\hat{\gamma}, \mathbf{w}, b} \frac{\hat{\gamma}}{\|\mathbf{w}\|}, \text{ introducing a functional margin as the relation } \gamma = \frac{\hat{\gamma}}{\|\mathbf{w}\|}$$
$$\text{s.t. } y^{(i)} \left( \mathbf{w}^T x^{(i)} + b \right) \geq \hat{\gamma}, i = 1, \dots, m$$

Now, we can freely choose  $\hat{\gamma}$  which requires us to find appropriate weights. We set  $\hat{\gamma} = 1$  and can now instead solve:

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2, \text{ s.t. } y^{(i)} \left( \mathbf{w}^T x^{(i)} + b \right) \geq 1, i = 1, \dots, m$$

# In General: Lagrangian Multiplier

We can reformulate a (primal) optimization problem:

$$\begin{aligned} \min_{\mathbf{w}} \quad & f(\mathbf{w}) \\ \text{s.t.} \quad & g_i(\mathbf{w}) \leq 0, i = 1, \dots, k \\ & h_j(\mathbf{w}) = 0, j = 1, \dots, l \end{aligned}$$

And instead solve the *generalized Lagrangian* (when the Karush-Kuhn-Tucker conditions are met):

$$\mathcal{L}(\mathbf{w}, \boldsymbol{\alpha}, \boldsymbol{\beta}) = f(\mathbf{w}) + \sum_{i=1}^k \alpha_i g_i(\mathbf{w}) + \sum_{j=1}^L \beta_j h_j(\mathbf{w})$$

$\alpha_i$  and  $\beta_j$  are the Lagrangian multipliers.

# Apply Lagrangian Multiplier

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2, \text{ s.t. } y^{(i)} \left( \mathbf{w}^T x^{(i)} + b \right) \geq 1, i = 1, \dots, m$$

We can formulate  $f(\mathbf{w}) = 1/2 \|\mathbf{w}\|^2$  and the constraints as:

$$g_i(\mathbf{w}) = -y^{(i)} \left( \mathbf{w}^T x^{(i)} + b \right) + 1 \leq 0, i = 1, \dots, m$$

and now solve the Lagrangian

$$\mathcal{L}(\mathbf{w}, b, \boldsymbol{\alpha}) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^m \alpha_i \left( y^{(i)} \left( \mathbf{w}^T x^{(i)} + b \right) - 1 \right)$$

by setting the derivatives to zero ( $\nabla_b \mathcal{L} = 0$  because of the KKT conditions):

$$\nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w}, b, \boldsymbol{\alpha}) = \mathbf{w} - \sum_{i=1}^m \alpha_i y^{(i)} x^{(i)} = 0 \Rightarrow \mathbf{w} = \sum_{i=1}^m \alpha_i y^{(i)} x^{(i)}$$

# Finding optimal weights

This gives us the optimal weights when having found the Lagrangian multipliers on our inputs:

$$\mathbf{w}^* = \sum_{i=1}^m \alpha_i y^{(i)} \mathbf{x}^{(i)}$$

Importantly, when now use this for a novel input  $\mathbf{x}'$  we would apply  $\mathbf{w}^{*T} \mathbf{x}' + b$  and decide classification based on the sign.

Using the above equality, we can reformulate this as

$$\begin{aligned} \mathbf{w}^{*T} \mathbf{x}' + b &= \left( \sum_{i=1}^m \alpha_i y^{(i)} \mathbf{x}^{(i)} \right)^T \mathbf{x}' + b \\ &= \sum_{i=1}^m \alpha_i y^{(i)} \langle \mathbf{x}^{(i)}, \mathbf{x}' \rangle + b \end{aligned}$$




# Largest margin Separation

- ▶ this only involves some data points (support vectors)
- ▶ the constrained optimization can be solved through a Lagrange multiplier
- ▶ this leads to the hyperplane decision function

$$\alpha_i \geq 0,$$

$$f(\mathbf{x}) = \text{sgn}\left(\sum_{i=1}^m \alpha^{(i)} y^{(i)} \langle \mathbf{x}, \mathbf{x}^{(i)} \rangle + b\right)$$




$$\max_{\mathbf{w}, b} \min\{\|\mathbf{x} - \mathbf{x}^{(i)}\|\}$$

with  $\langle \mathbf{w}, \mathbf{x} \rangle + b = 0$  defining the hyperplane

# **Application of Kernel**

# Support Vector Machine



SVMs go back to (Vapnik 1998) , and a good tutorial can be found in (Burges 1998).

# Kernel Trick

The kernel trick for kernel methods as SVMs is a substitution:

- ▶ All computations can be formulated in a scalar product space.
- ▶ We introduce a kernel function – this express the scalar product in the higher dimensional feature space in terms of the lower-dimensional input space.
- ▶ The kernel function evaluates the function and scalar product of the feature space only from the lower-dimensional input space.

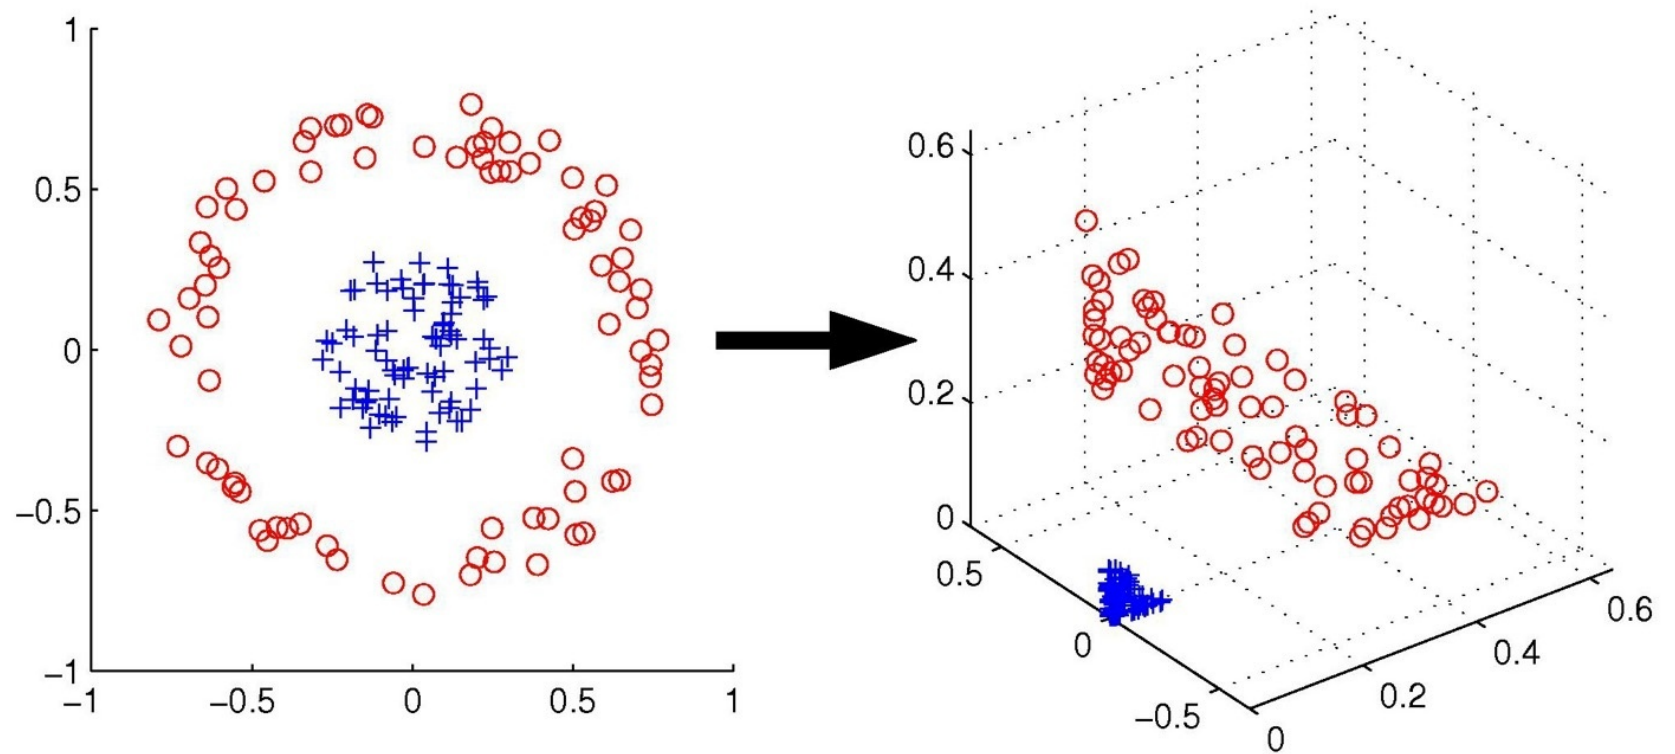
$$\text{e.g., } k(x, x') := \langle \mathbf{x}, \mathbf{x}' \rangle$$

# Recap – Example for Application of Kernel

Kernel functions provide mappings that allow for separability:

$$\phi(\mathbf{x}) \rightarrow x_1^2, x_2^2, \sqrt{2}x_1x_2$$

Importantly, the scalar product is not computed explicitly in the feature space. It can be applied in the input space – Kernel Trick.




# Kernel trick

$$\phi(\mathbf{x}) \rightarrow x_1^2, x_2^2, \sqrt{2}x_1x_2$$

$$\begin{aligned}\langle \phi(\mathbf{x}), \phi(\mathbf{x}') \rangle &= \langle (x_1^2, x_2^2, \sqrt{2}x_1x_2), (x_1'^2, x_2'^2, \sqrt{2}x_1'x_2') \rangle \\ &= \underbrace{x_1^2x_1'^2}_{a^2} + \underbrace{x_2^2x_2'^2}_{b^2} + \underbrace{2x_1x_2x_1'x_2'}_{2ab} \\ &= \left( \underbrace{x_1x_1'}_a + \underbrace{x_2x_2'}_b \right)^2 \\ &= \langle \mathbf{x}, \mathbf{x}' \rangle^2 = k(\mathbf{x}, \mathbf{x}')\end{aligned}$$

# Summary: Support Vector Machine



- ▶ Support vector machines implement the large margin principle.
- ▶ They apply non-linear mappings.
- ▶ Importantly, the scalar product is not computed explicitly in the feature space. using the Kernel Trick. This is much more efficient.
- ▶ The kernel function (weighted by multipliers) is applied wrt. the support vectors.

SVMs go back to (Vapnik 1998) , and a good tutorial can be found in (Burges 1998).

# Kernel functions

Polynomial kernel (of degree  $d$ )

$$k(x, x') := \langle x, x' \rangle^d$$

- Includes all polynomial terms up to degree  $d$ .

Radial Basis Function kernels

$$k(x, x') := \exp\left(-\frac{\|x - x'\|^2}{2\sigma^2}\right)$$

- Allows to map into an infinite dimensional feature space (Gaussian kernel can be constructed from an infinite sum over polynomial kernels).
- Scales with number of features.



# SVM – Advantages

- ▶ Very robust, guaranteed to be a global minimum
- ▶ Work well on small (and high dimensional) data spaces.
- ▶ Does allow for non-linearly separable data (using Kernel trick).
- ▶ Can be softened through a simple parameter allowing for violation of the maximum margin.
- ▶ Is efficient for high-dimensional datasets as the complexity is characterized by the number of support vectors.
- ▶ Support Vectors can help to understand the problem better.
- ▶ Only a small number of hyperparameters.

# **SVM – Disadvantages**

- ▶ Not suitable for big datasets as the training time with SVMs becomes much more computationally intensive.
- ▶ They are less effective on noisier datasets with overlapping classes.
- ▶ Are often outperformed by Deep Neural Networks.

**Commercial Break**

# CITEC Conference today and tomorrow



For information see [Conference Webpage](#).



# Reservoir Computing

# From Dynamical Features ...

Temporal Filters can be seen as dynamical systems that compute at each time step a state that is some function of previous states and the current input:

$$s_t = F(s_{t-1}, s_{t-2}, \dots, s_{t-k}, x_t, x_{t-1}, \dots, x_{t-l})$$

In the simplest case, the function is linear in its arguments, leading to the well-known recursive filters

$$s_t = \sum_{i=1}^K a_i s_{t-i} + \sum_{j=0}^L b_j x_{t-j}$$

that allow, e.g., to selectively damp/enhance specifiable frequency bands of the input time sequence. For example a smoothing filter:

$$s_t = (1 - \gamma)s_{t-1} + \gamma x_t$$

## **... to Dynamical Systems**

Yet, combining linear filters always leads back to a linear filter.

Richer processing can only occur when non-linearities are included.

For example, we can consider non-linear filters arising from recurrent neural networks as in reservoir computing.



# Learning from Random Features

Simple learning approach in a feedforward neural network:

- ▶ using randomly initialized early layers and keep them fixed (comparable to expansion in SVMs) – use large input layers that provide diversity
- ▶ During learning: only adapt output weights – linear transformation of the (random) features.
- ▶ Such an expansion of the input space can facilitate learning and allow for better separability.

# Random Features in a Recurrent Neural Net

Echo state networks apply the same idea in a recurrent neural network:

- ▶ Initialize the recurrent neural network randomly and keep it fixed.
- ▶ The same holds true for the projection of the input onto the recurrent layers.
- ▶ Only train the connections towards the output layer which makes learning very simple.
- ▶ The recurrent part is called a reservoir – it should cover a diversity of dynamics that can be recruited.

Following (Hinton 2013) in his Advanced Machine Learning Course.

# Setting the Connections inside the Reservoir

Crucial for Echo State Networks is the setup of the random recurrent connections:


- ▶ They have to be kept bound and fulfill the echo state property to prevent dying or exploding activations.
- ▶ Still, activities may decay too fast or too slowly. Therefore, the reservoir has to be tuned in a way that the dynamics of the features match the time scales of the application task.

## Echo Property


Without external excitation all activities of the reservoir will decay slowly to zero. A criterion for this is that the spectral radius (the largest eigenvalue of  $A^T A$ ) is less than 1 (or set to 1).

# Echo State Network

- ▶ Input projects onto reservoir, here a real value.
- ▶ Target output: is a sine wave with the frequency given by the frequency given by the input.



# Echo State Network Example Results



A test run of the frequency generator from the previous slide.

In the back, the input step function is shown.

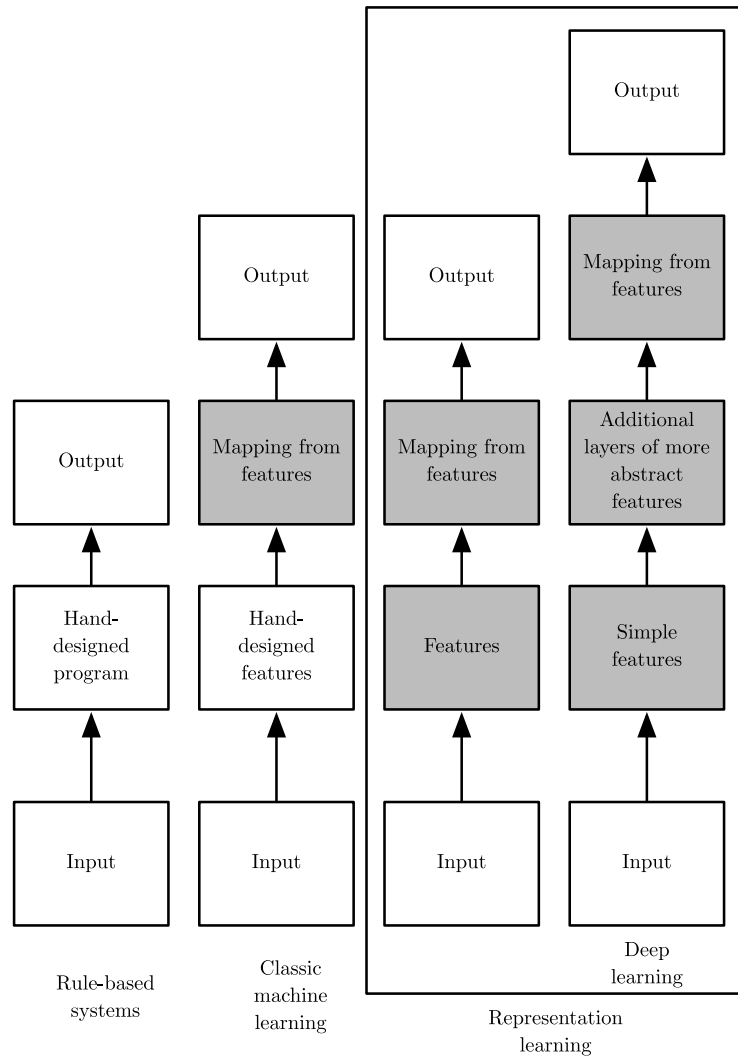
The black sinewaves is the target output (unknown to the network).



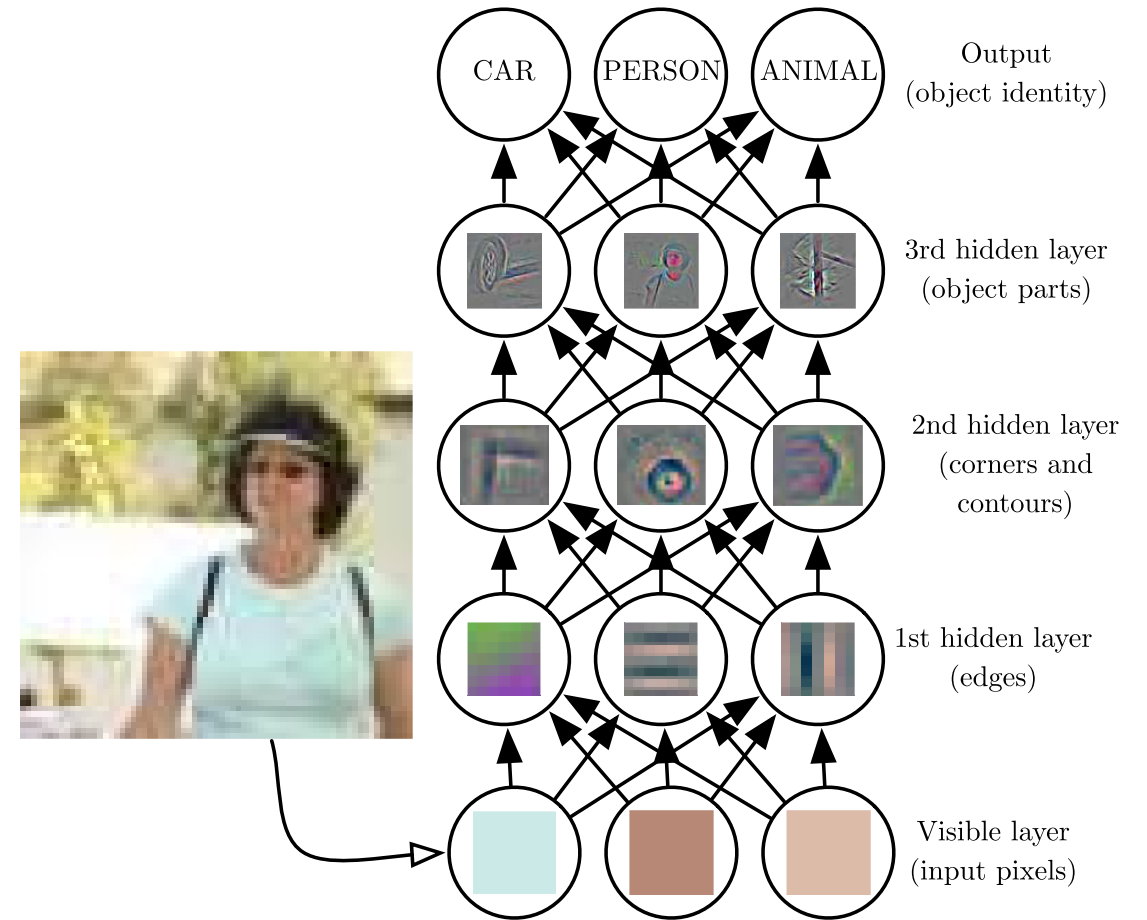
# Recap – Representation Learning

# Representation Learning

## Current ML Pipeline



## End-to-End Learning in Deep NN



(Goodfellow, Bengio, and Courville 2016)



# Example: Waymo


Scene Representation in Autonomous Driving



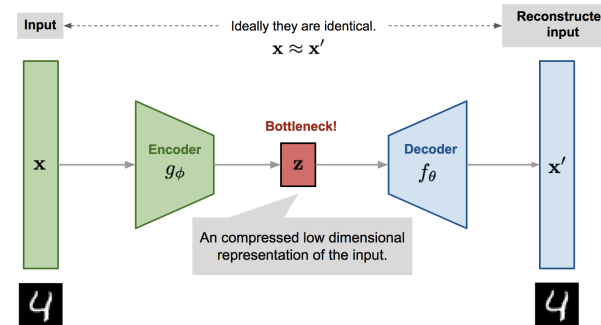
**Features: Transfer Learning**

# Autoencoder (Weng 2018)


## Entangled Representation



## Autoencoder




## Disentangled Representation



- ▶ Encoder translates high-dimension input into latent low-dimensional code.
- ▶ Decoder recovers data from the code.

# Support Vector Machine



- ▶ Support vector machines implement the large margin principle.
- ▶ They apply non-linear mappings.
- ▶ Importantly, the scalar product is not computed explicitly in the feature space. using the Kernel Trick. This is much more efficient.
- ▶ The kernel function (weighted by multipliers) is applied wrt. the support vectors.

SVMs go back to (Vapnik 1998) , and a good tutorial can be found in (Burges 1998).

# References

- Burges, Christopher J. C. 1998. “A Tutorial on Support Vector Machines for Pattern Recognition.” *Data Mining and Knowledge Discovery* 2: 121–67.
- Goodfellow, Ian, Yoshua Bengio, and Aaron Courville. 2016. *Deep Learning*. MIT Press.
- Hinton, Geoffrey E. 2013. “Recurrent Neural Networks.” CSC 2535: Advanced Machine Learning Course.
- Jaeger, Herbert. 2007. “Echo State Network.” *Scholarpedia* 2 (9): 2330. <https://doi.org/10.4249/scholarpedia.2330>.
- Nayak, Sunita. 2019. “Image Classification Using Transfer Learning in Pytorch.” 2019. <https://www.learnopencv.com/image-classification-using-transfer-learning-in-pytorch/>.
- Ng, Andrew. 2018. “Support Vector Machines.” Course CS229, Stanford University, Lecture Notes.
- Vapnik, Vladimir N. 1998. *Statistical Learning Theory*. Wiley-Interscience.

